



Universitat
Pompeu Fabra
Barcelona

Globally Optimal Hierarchical Reinforcement Learning for Linearly-Solvable Markov Decision Processes

Guillermo Infante, Anders Jonsson, Vicenç Gómez

AAAI 2022

Overview

Introduction

Background

Hierarchical LMDPs

Algorithms

Experiments and results

Contributions and Conclusion

Introduction

Introduction

- Hierarchical Reinforcement Learning aims to make **learning more efficient** by decomposing large problems.
- Novel approach to HRL using Linearly-solvable Markov Decision Processes (LMDPs).
- LMDPs are **computationally efficient** (Todorov, 2006).
- **Globally optimal value function** (vs. *hierarchically optimal* or *recursively optimal*) (Dietterich, 2000).

Background

Background - LMDPs (i)

An LMDP (Kappen et al., 2012; Todorov, 2006) is as a tuple $\mathcal{L} = \langle \mathcal{S}, \mathcal{T}, \mathcal{P}, \mathcal{R}, \mathcal{J} \rangle$:

- We define $\mathcal{S}^+ = \mathcal{S} \cup \mathcal{T}$ to denote the full set of states.
- $\mathcal{P} : \mathcal{S} \rightarrow \Delta(\mathcal{S}^+)$.
- $\mathcal{R} : \mathcal{S} \rightarrow \mathbb{R}$ and $\mathcal{J} : \mathcal{T} \rightarrow \mathbb{R}$.
- Learning agent follows a policy $\pi : \mathcal{S} \rightarrow \Delta(\mathcal{S}^+)$.

Background - LMDPs (ii)

- At time-step t , the agent observes s_t and receives a reward

$$\mathcal{R}(s_t, \pi) = \mathcal{R}(s_t) - \lambda \cdot \text{KL}(\pi(\cdot|s_t) \| \mathcal{P}(\cdot|s_t)).$$

- The aim of the agent is to maximize

$$v^\pi(s) = \mathbb{E} \left[\sum_{t=1}^{T-1} \mathcal{R}(S_t, \pi) + \mathcal{J}(S_T) \mid S_1 = s \right].$$

- We obtain the following Bellman optimality equation

$$\frac{1}{\lambda} v(s) = \frac{1}{\lambda} \mathcal{R}(s) + \max_{\pi} \mathbb{E}_{s' \sim \pi(\cdot|s)} \left[\frac{1}{\lambda} v(s') - \log \frac{\pi(s'|s)}{\mathcal{P}(s'|s)} \right] \quad (\forall s).$$

Background - LMDPs (iii) - Eigenvector setting

- Taking $z(s) = e^{v(s)/\lambda}$ for each $s \in \mathcal{S}^+$ leads to

$$z(s) = e^{\mathcal{R}(s)/\lambda} \sum_{s'} \mathcal{P}(s'|s) z(s').$$

- If \mathcal{P} and \mathcal{R} are known, then

$$\mathbf{z} = R P \mathbf{z}^+ \text{ where } R = \text{diag}(e^{\mathcal{R}(\cdot)/\lambda}).$$

Background - LMDPs (iii) - Online setting

- Alternatively, there is an online (corrected) update rule

$$\hat{z}(s_t) \leftarrow (\mathbf{1} - \alpha_t) \hat{z}(s_t) + \alpha_t e^{r_t/\lambda} \hat{z}(s_{t+1}) \frac{\mathcal{P}(s_{t+1}|s_t)}{\hat{\pi}(s_{t+1}|s_t)},$$

samples are in the form of (s_t, r_t, s_{t+1}) .

- Given a z , policies are derived using

$$\pi(s'|s) = \frac{\mathcal{P}(s'|s)z(s')}{\sum_{s''} \mathcal{P}(s''|s)z(s'')}.$$

Background - LMDPs (iv) - Compositionality

- Let $\{\mathcal{L}_1, \dots, \mathcal{L}_n\}$ be a collection of LMDPs.
- Each LMDP \mathcal{L}_i only differs in $\mathcal{J}_i(t)$.
- For a new LMDP \mathcal{L} for which the next holds (Todorov, 2009)

$$e^{\mathcal{J}(t)/\lambda} = z(t) = \sum_{k=1}^n w_k e^{\mathcal{J}_k(t)/\lambda} \text{ for } t \in \mathcal{T}_i.$$

- Due to linearity, the following is also satisfied

$$z(s) = \sum_{k=1}^n w_k z_k(s) \quad \forall s \in \mathcal{S}.$$

Hierarchical LMDPs

Hierarchical LMDPs (i)

- Inspired by Wen et al. (Wen et al., 2020).
- For an LMDP \mathcal{L} , its \mathcal{S} is partitioned into L subsets $\{\mathcal{S}_i\}_{i=1}^L$.
- Each \mathcal{S}_i induces a subtask $\mathcal{L}_i = \langle \mathcal{S}_i, \mathcal{T}_i, \mathcal{P}_i, \mathcal{R}_i, \mathcal{J}_i \rangle$:
 - ▶ \mathcal{T}_i includes states in $\mathcal{S}^+ \setminus \mathcal{S}_i$ that are one step away from any $s \in \mathcal{S}_i$.
 - ▶ For $\tau \in \mathcal{T}_i$,

$$\mathcal{J}_i(\tau) = \begin{cases} \mathcal{J}(\tau) & \text{if } \tau \in \mathcal{T}, \\ \hat{v}(\tau) & \text{otherwise.} \end{cases}$$

Hierarchical LMDPs (ii)

Definition

Subtask equivalence (for some \mathcal{L}_i and \mathcal{L}_j) implies a bijective relationship $f : \mathcal{S}_i \rightarrow \mathcal{S}_j$.

- Set of *exit states* $\mathcal{E} = \cup_{i=1}^L \mathcal{T}_i$
- $\mathcal{E}_i = \mathcal{E} \cap \mathcal{S}_i$ denotes the set of *exit states inside* subtask \mathcal{L}_i .
- Set of *equivalence classes* $\mathcal{C} = \{\mathcal{C}_1, \dots, \mathcal{C}_C\}$, $C \leq L$.
- A single subtask $\mathcal{L}_j = \langle \mathcal{S}_j, \mathcal{T}_j, \mathcal{P}_j, \mathcal{R}_j, \mathcal{J}_j \rangle$ per equivalence class $\mathcal{C}_j \in \mathcal{C}$.

Hierarchical LMDPs (iii) - Illustration




Figure 1: a) A 4-room LMDP, with all exit states highlighted; b) a single subtask with 5 terminal states G, L, R, T, B that is equivalent to all 4 room subtasks.

Hierarchical LMDPs (iv) - Subtask compositionality (i)

- Consider subtask \mathcal{L}_j and its terminal set $\mathcal{T}_j = \{\tau_1, \dots, \tau_n\}$.
- We define n base LMDPs $\mathcal{L}_j^1, \dots, \mathcal{L}_j^n$, which only differ in \mathcal{J}_j^k .
- Concretely,

$$z_j^k(\tau) = \begin{cases} 1 & \text{if } \tau = \tau_k \rightarrow \mathcal{J}_j^k(\tau) = 0, \\ 0 & \text{otherwise} \rightarrow \mathcal{J}_j^k(\tau) = -\infty \end{cases}$$

- Solutions are optimal z_j^1, \dots, z_j^n .
- Having $\hat{z}(s)$ for each $t \in \mathcal{T}_j$, compositionality rule

$$\hat{z}(s) = \hat{z}(\tau_1)z_j^1(s) + \dots + \hat{z}(\tau_n)z_j^n(s) \quad \forall s \in \mathcal{S}_i, \forall \mathcal{L}_i \in \mathcal{C}_j.$$

Hierarchical LMDPs (v) - Subtask compositionality (ii)

- For all subtasks, terminal states $\tau_1 \dots \tau_n$ are by definition in \mathcal{E} .
- Enough represent $\hat{z}(s) \forall s \in \mathcal{S}$:
 - ▶ $\hat{z}_{\mathcal{E}} : \mathcal{E} \rightarrow \mathbb{R}$,
 - ▶ for the base LMDPs z_j^1, \dots, z_j^n .
- Solutions for base LMDPs can be reused.

Algorithms

Eigenvector algorithm (i)

- We can restrict

$$\hat{z}(s) = \hat{z}(\tau_1)z_j^1(s) + \cdots + \hat{z}(\tau_n)z_j^n(s) \quad \forall s \in \mathcal{S}_i, \forall \mathcal{L}_i \in \mathcal{C}_j.$$

to states $s \in \mathcal{E}$.

- Thus, we define

$$\mathbf{z}_{\mathcal{E}} = G\mathbf{z}_{\mathcal{E}}.$$

- This corresponds to an eigenvector problem.
- Value estimates for $s \in \mathcal{S} \setminus \mathcal{E}$ are obtained afterwards.

Online algorithm (i)

- We keep $\hat{z}_j^1, \dots, \hat{z}_j^n$ for each \mathcal{L}_j .
- We can update all base LMDPs within a subtask \mathcal{L}_j using intra-task learning (Kaelbling, 1993; Jonsson and Gómez, 2016).
- The online update rule (again, restricted to $s \in \mathcal{E}$)

$$\hat{z}_{\mathcal{E}}(s) \leftarrow (1 - \alpha_{\ell})\hat{z}_{\mathcal{E}}(s) + \alpha_{\ell}[\hat{z}_{\mathcal{E}}(t_1)\hat{z}_j^1(s) + \dots + \hat{z}_{\mathcal{E}}(t_n)\hat{z}_j^n(s)].$$

- Estimates at any level are learned in an episodic manner.
- When shall we update states in \mathcal{E} ?

Online algorithm (ii)

We propose the following alternatives:

V_1 : Update $s \in \mathcal{E}_i$ each time the agent transitions from s .

V_2 : When the agent reaches $\tau \in \mathcal{T}_i$ of the subtask \mathcal{L}_i , update the values of every $s \in \mathcal{E}_i$.

V_3 : When the agent reaches $\tau \in \mathcal{T}_i$ of the subtask \mathcal{L}_i , update the values of every $s \in \mathcal{E}_i$ and all exit states of subtasks in the equivalence class \mathcal{C}_j of \mathcal{L}_i .

Experiments and results

Results - Rooms domain




Figure 2: MAE over time for 3×3 (top-left), 5×5 (top-right) and 8×8 (bottom) room instances.

Results - Taxi domain




Figure 3: MAE over time for 5×5 (left) and 10×10 (right) grids of Taxi domain.

Contributions and Conclusion

Contributions

- Novel scheme based on subtask compositionality.
- The subtasks decomposition is at the level of the value function.
- Our method converges to the optimal value function.

Conclusion

- We are able to retrieve the optimal value function.
- New form of zero-shot learning.

References

- Dietterich, T. G. (2000). Hierarchical reinforcement learning with the MAXQ value function decomposition. *J. Artif. Intell. Res.*, 13:227–303.
- Jonsson, A. and Gómez, V. (2016). Hierarchical Linearly-Solvable Markov Decision Problems. In *Proceedings of the 26th International Conference on Automated Planning and Scheduling (ICAPS)*.
- Kaelbling, L. P. (1993). Learning to Achieve Goals. In *Proceedings of the International Joint Conference on Artificial Intelligence (IJCAI)*, pages 1094–1099.

- Kappen, H. J., Gómez, V., and Opper, M. (2012). Optimal control as a graphical model inference problem. *Mach. Learn.*, 87(2):159–182.
- Todorov, E. (2006). Linearly-solvable Markov decision problems. *Advances in Neural Information Processing Systems (NIPS)*, pages 1369–1376.
- Todorov, E. (2009). Compositionality of optimal control laws. *Advances in Neural Information Processing Systems (NIPS)*, pages 1856–1864.
- Wen, Z., Precup, D., Ibrahimi, M., Barreto, A., Van Roy, B., and Singh, S. (2020). On Efficiency in Hierarchical Reinforcement Learning. In *Proceedings of the 34th Conference on Neural Information Processing Systems (NeurIPS)*.